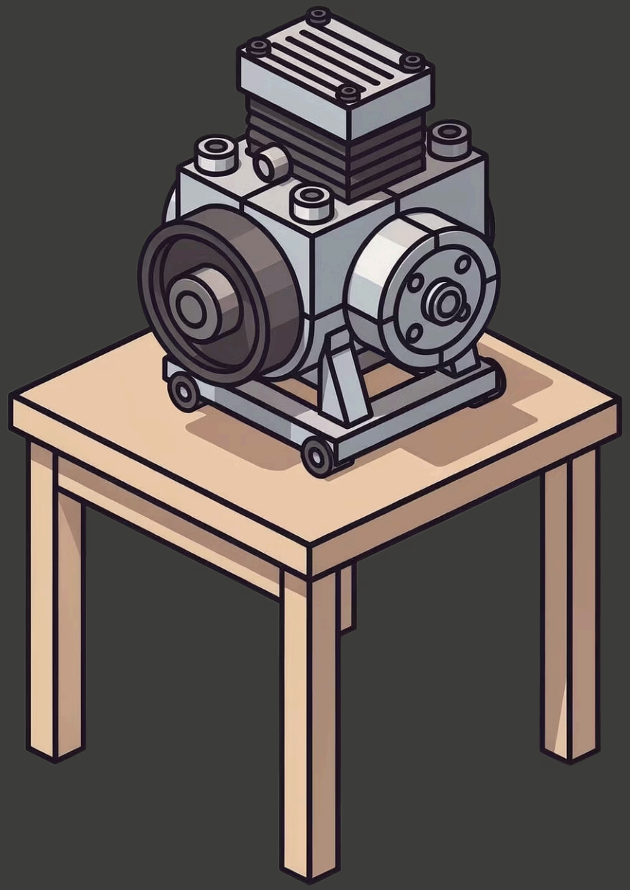

GARAGE_INFERENCE

Big Ideas. Cheap Models. The constraint is the creativity.

MAY 1–4, 2026

ONLINE · FREE

$1,700 IN PRIZES

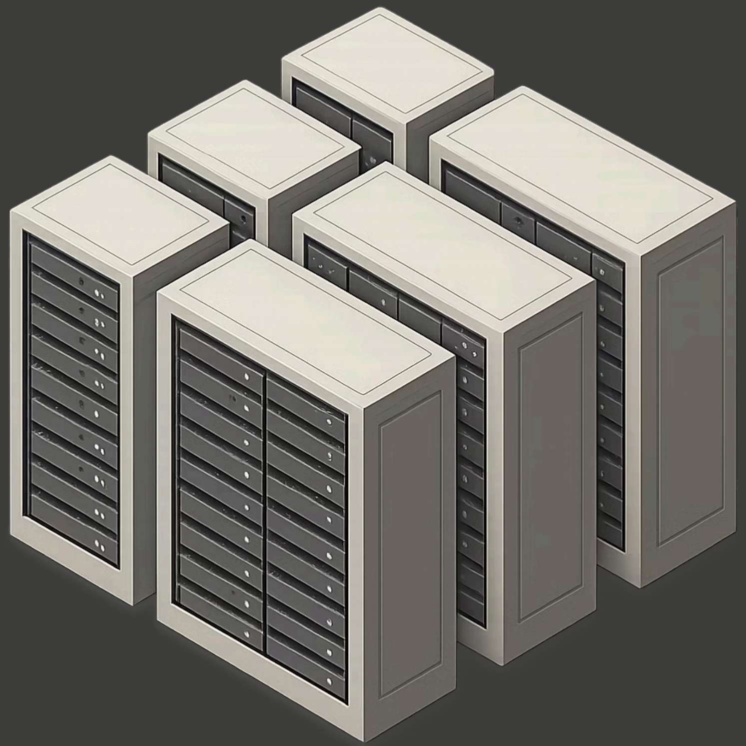

THE PROBLEM

Most hackathons reward the biggest model, the fattest GPU, the most expensive API call. That's not where the real engineering lives.

Rate-Limited & Expensive

Frontier models are rate-limited and expensive — they're not accessible to most builders.

Budget Exclusion

$200/month API budgets exclude billions of builders from participating in the AI revolution.

Offline Blind Spot

Cloud endpoints don't reach offline devices. Entire use cases are left on the table.

Not a Craft

"Just throw GPT-5 at it" is not engineering. It's a shortcut that skips the hard thinking.

> The future of AI isn't the biggest model — it's the **smallest model that still gets the job done.**

YOUR MISSION

Build a working product powered by the **cheapest, weakest LLMs you can find** — local or budget-tier API. Then show the gap between how bad your model is and how good your project turned out.

best_project = weakest_possible_model × highest_possible_impact

A solid app on Gemma 3 1B beats a decent app on GPT-5 Nano. Every time.

ONE REAL PROBLEM

ONE CHEAP MODEL

72 HOURS

The Rules of the Game

- Use the weakest model you can justify
- Engineer around its limitations
- Show your work — document everything
- Prove the gap between model quality and project quality
- Open-source everything

4 MODEL TIERS — PICK YOUR WEIGHT CLASS

Lower tier = more points. Same output quality on a weaker model is harder — judges reward it.

Tier 1 · Absolute Garage MAX

Barely follows instructions

- Qwen 3 0.6B / 1.7B
- Gemma 3 1B, Gemma 3n E2B
- SmolLM2 135M–1.7B
- Phi-4 Mini, Ministral 3B
- Browser / edge models
- Anything ≤4B params

Tier 2 · Proper Garage HIGH

Cheap but reliable

- Qwen 3 4B / 8B
- Gemma 3 4B (Q4), Gemma 3n E4B
- Ministral 8B
- DeepSeek-R1-Distill 7B / 8B
- API: GPT-5 Nano
- Gemini 2.0 Flash-Lite / Flash

The lower your tier, the more the judges lean in. Tier 1 projects that deliver real results are the ones people remember.

TIERS — CONTINUED

Tier 3 · Home Workshop MOD

Decent at low prices

- Qwen 3 14B / 32B
- Qwen 3 30B-A3B (MoE)
- Gemma 3 12B / 27B (Q4)
- Ministral 14B
- API: GPT-5 Mini
- Gemini 2.5 Flash ($0.25–$1.50/M tokens)

Tier 4 · Rented Studio LOW

Real capability — prove the engineering mattered

- Llama 4 Scout (Q4+)
- DeepSeek-R1-Distill 70B
- Qwen 3 235B-A22B (Q4)
- API: Claude Haiku 4.5
- Gemini 2.5 Pro batch ($1.50–$3/M)

⊗ 🚫 **BANNED:** GPT-5/5.2/5.2 Pro, Claude Sonnet 4.6/Opus 4.6, Gemini 3 Pro, o3-pro, Llama 4 Maverick (full), anything >$3/M input tokens, any frontier model. When in doubt — drop a tier. Judges reward it.

WHAT A GOOD SUBMISSION LOOKS LIKE

Your model is dumb. Your job is to engineer around its weaknesses. Pick the techniques that fit your project:

Structured Prompts

Constrain output to what the model can handle. Shape the input, shape the result.

RAG

Model doesn't need to know — just read and respond. Retrieval does the heavy lifting.

Tool Use

Orchestrate actions, don't generate raw knowledge. Let tools do what models can't.

Multi-Step Pipelines

Break tasks into tiny subtasks. Small models excel at narrow, well-defined jobs.

Validation Layers

Catch bad outputs and retry. Build the safety net your model needs.

Caching & Fine-tuning

Pre-compute what you can. Make a tiny model an expert at one specific thing.

Show the judges you thought about it. Not a benchmark wrapper. Not a chatbot that just chats.

WHAT YOU'RE SHIPPING

Every submission must include all six of the following components. No exceptions.

1

Working Demo

Live URL, video, or installable project. **Must run.**

2

Public GitHub Repo

Clean README, one-command setup. No mystery installs.

3

Exact Model Declaration

"Qwen 3 4B Q4_K_M" — not "a small Llama." Precision matters.

4

Technical Writeup

What model, what problem, what techniques, honest division of labor.

5

Cost & Performance Metrics

Actual $ spent, latency, tokens/sec. Show the receipts.

6

2-Min Demo Video + Known Failures

Where does the model still break? **Honesty is rewarded.**

All winning projects must be open-source.

JUDGING CRITERIA — PART 1

⭐ The Wow Gap

30%

The core metric. How wide is the gap between your model's expected capability and what your project actually does? A Tier 1 model producing Tier 3 results wins. Every time.

Practical Usefulness

20%

Would a real person install this, bookmark it, pay for it? Not "would they be impressed at demo day" — would they actually use it next week?

Technical Execution

20%

Stable, fast, smart trade-offs, graceful failures, clean code. Smooth UX. Thoughtful architecture documented in the writeup.

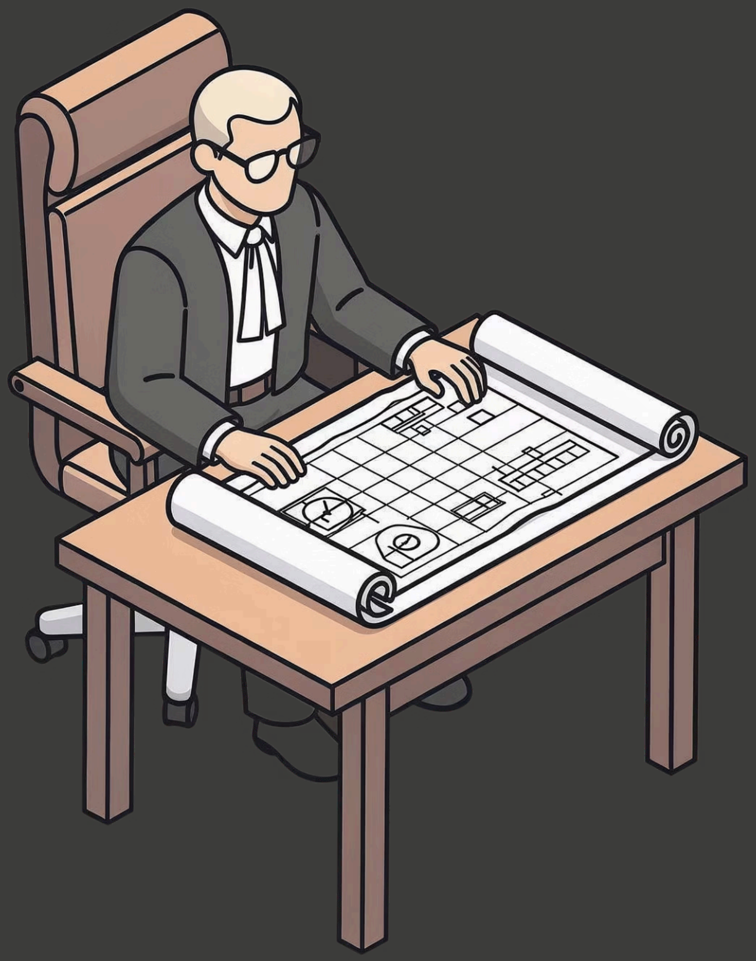

JUDGING CRITERIA — PART 2

Creativity & Innovation

10%

Novel angle. Unexpected use case. A technique judges haven't seen. Not another standard RAG chatbot with fresh paint.

Accessibility & Reproducibility

10%

One-command setup. Clear README. Runs on a $500 laptop. Could a student in a developing country use this?

Secure Design

10%

Small models hallucinate, jailbreak, follow injected instructions. Least-privilege scopes. Input/output validation. No raw shell or DB access from model output. Clear threat model.

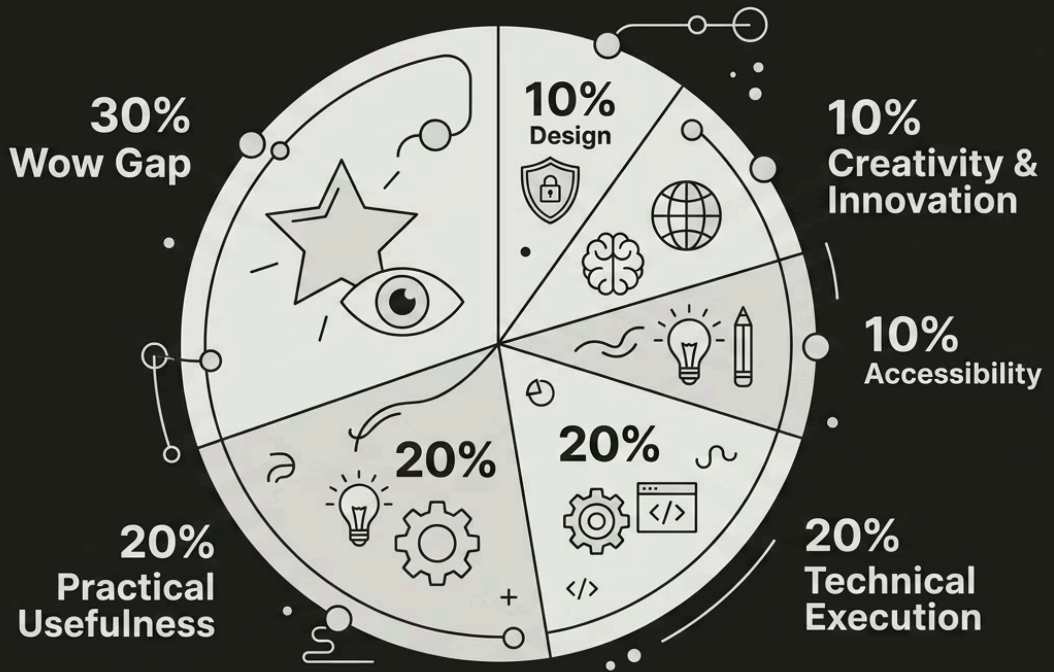

The tier system is a lens, not a multiplier. Same score, lower tier wins.

PRIZES — $1,700

$1,000

🥇 Root Access Award

Best weak model + strong engineering combination. The ultimate Garage_Inference prize.

$300

🥈 Runner-Up

Outstanding execution and creative squeeze from a constrained model.

$200

🥉 Finalist

Originality with the least capable model. Proof that constraints breed creativity.

$200

♡ Community Choice

Voted by participants. The project the builders loved most.

Beyond the Cash

- Code review from senior engineers
- Personalized feedback on your project
- Professional connections in the AI space
- A portfolio-worthy shipped project

Payment Details

Prizes paid via **PayPal or bank transfer** within 14 days of winner announcement. No hoops, no delays.

✓ All prizes guaranteed. No sponsorship dependencies.

TIMELINE

1
May 1 · 18:00 UTC

Hackathon starts. **72 hours on the clock.** Build begins.

2
May 4 · 18:00 UTC

Submissions close. Push your repo and submit. No late entries.

3
May 4 – May 16

Judging panel reviews every submission. Every project gets a real look.

4
May 16 – May 23

Community voting opens. One project, one vote. Make your project shareable.

5
May 23

Winners announced. Prizes paid within 14 days of announcement.

BEFORE YOU START — RIGHT NOW

Answer these four questions before you write a single line of code:

1 What's my model?

Exact name + tier. **"Qwen 3 0.6B"**, not "small one."
Precision is the first sign of good engineering.

2 What real problem am I solving?

Someone uses it next week — yes or no? If you can't answer yes, find a different problem.

3 What's my engineering scaffolding?

RAG, structured outputs, tool use, validation — pick 2–3.
Know your stack before you start.

4 Where does inference run?

Browser, CPU laptop, RPi, cheap API — and what's it cost?
Know your constraints cold.

If you answered all four — you're ready.

> **72 hours. The weakest model you can find. The smartest engineering you can muster.**
> Build something that makes people forget the model is cheap.